

February 1, 2024

## 1 Logistic Regression In R

```
[21]: # Summary of dataset in package
summary(mtcars)
```

mpg	cyl	disp	hp
Min. :10.40	Min. :4.000	Min. : 71.1	Min. : 52.0
1st Qu.:15.43	1st Qu.:4.000	1st Qu.:120.8	1st Qu.: 96.5
Median :19.20	Median :6.000	Median :196.3	Median :123.0
Mean :20.09	Mean :6.188	Mean :230.7	Mean :146.7
3rd Qu.:22.80	3rd Qu.:8.000	3rd Qu.:326.0	3rd Qu.:180.0
Max. :33.90	Max. :8.000	Max. :472.0	Max. :335.0

  

drat	wt	qsec	vs
Min. :2.760	Min. :1.513	Min. :14.50	Min. :0.0000
1st Qu.:3.080	1st Qu.:2.581	1st Qu.:16.89	1st Qu.:0.0000
Median :3.695	Median :3.325	Median :17.71	Median :0.0000
Mean :3.597	Mean :3.217	Mean :17.85	Mean :0.4375
3rd Qu.:3.920	3rd Qu.:3.610	3rd Qu.:18.90	3rd Qu.:1.0000
Max. :4.930	Max. :5.424	Max. :22.90	Max. :1.0000

  

am	gear	carb
Min. :0.0000	Min. :3.000	Min. :1.000
1st Qu.:0.0000	1st Qu.:3.000	1st Qu.:2.000
Median :0.0000	Median :4.000	Median :2.000
Mean :0.4062	Mean :3.688	Mean :2.812
3rd Qu.:1.0000	3rd Qu.:4.000	3rd Qu.:4.000
Max. :1.0000	Max. :5.000	Max. :8.000

Logistic regression is implemented in R using `glm()` by training the model using features or variables in the dataset.

```
[22]: # Installing the package

# For Logistic regression
install.packages("caTools")

# For ROC curve to evaluate model
install.packages("ROCR")
```

```
# Loading package
library(caTools)
library(ROCR)
```

Installing package into ‘/usr/local/lib/R/site-library’  
(as ‘lib’ is unspecified)

Installing package into ‘/usr/local/lib/R/site-library’  
(as ‘lib’ is unspecified)

```
[23]: # Splitting dataset
split <- sample.split(mtcars, SplitRatio = 0.8)
split

train_reg <- subset(mtcars, split == "TRUE")
test_reg <- subset(mtcars, split == "FALSE")

# Training model
logistic_model <- glm(vs ~ wt + disp,
                      data = train_reg,
                      family = "binomial")

logistic_model

# Summary
summary(logistic_model)
```

1. FALSE 2. TRUE 3. TRUE 4. TRUE 5. TRUE 6. FALSE 7. TRUE 8. TRUE 9. TRUE 10. TRUE  
11. FALSE

Call: glm(formula = vs ~ wt + disp, family = "binomial", data = train\_reg)

Coefficients:

(Intercept)	wt	disp
1.44440	1.54534	-0.03262

Degrees of Freedom: 23 Total (i.e. Null); 21 Residual

Null Deviance: 33.1

Residual Deviance: 15.8 AIC: 21.8

Call:

glm(formula = vs ~ wt + disp, family = "binomial", data = train\_reg)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.44440	3.35050	0.431	0.6664
wt	1.54534	1.94984	0.793	0.4280

```
disp          -0.03262    0.01797  -1.815    0.0695 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 33.104  on 23  degrees of freedom
Residual deviance: 15.801  on 21  degrees of freedom
AIC: 21.801
```

Number of Fisher Scoring iterations: 6

## 1.1 Predict test data based on model

```
[24]: predict_reg <- predict(logistic_model,
                             test_reg, type = "response")
predict_reg
```

```
Mazda RX4 0.56811772067142 Valiant 0.366336432517721 Merc 280C 0.784730885831396
Merc 450SE 0.220569990126755 Chrysler Imperial 0.00948925420350044 Dodge Challenger
0.0296325009949015 AMC Javelin 0.0405629141446895 Lotus Europa 0.663838948193045
```

```
[26]: # Changing probabilities
predict_reg <- ifelse(predict_reg > 0.5, 1, 0)

# Evaluating model accuracy
# using confusion matrix
table(test_reg$vs, predict_reg)

missing_classerr <- mean(predict_reg != test_reg$vs)
print(paste('Accuracy =', 1 - missing_classerr))

# ROC-AUC Curve
ROCPred <- prediction(predict_reg, test_reg$vs)
ROCPer <- performance(ROCPred, measure = "tpr",
                      x.measure = "fpr")

auc <- performance(ROCPred, measure = "auc")
auc <- auc@y.values[[1]]
auc

# Plotting curve
plot(ROCPer)
plot(ROCPer, colorize = TRUE,
     print.cutoffs.at = seq(0.1, by = 0.1),
     main = "ROC CURVE")
abline(a = 0, b = 1)
```

```
auc <- round(auc, 4)
legend(.6, .4, auc, title = "AUC", cex = 1)
```

```
predict_reg
  0 1
0 4 1
1 1 2

[1] "Accuracy = 0.75"
0.733333333333333
```



